

OpenCL exercise 3: Sobel filter

Kaicong Sun

Sobel filter

- ▶ Used for edge detection in images
- ▶ A combination of two convolutional operators: horizontal and vertical

$$G_1 = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} * A$$


$$G_2 = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} * A$$

$$G = \sqrt{G_1^2 + G_2^2}$$


Sobel filter

```
1 G1(i,j)=
2     1*A(i-1, j-1) +2*A(i-1, j  ) +1*A(i-1, j+1)
3     +           0      +           0      +           0
4     -1*A(i+1, j-1) -2*A(i+1, j  ) -1*A(i+1, j+1)
5
6 G2(i,j)=
7     1*A(i-1, j-1) +           0      -1*A(i-1, j+1)
8     +2*A(i  , j-1) +           0      -2*A(i  , j+1)
9     +1*A(i+1, j-1) +           0      -1*A(i+1, j+1)
10
11 G(i,j)=sqrt(G1(i,j)*G1(i,j)+ G2(i,j)*G2(i,j))
```


Sobel filter




Original image



G_1




G_2




Output image

Sobel filter



Original image



Output of Sobel filter

Task 1

- ▶ Implement the sobel filter on the GPU, similar to the CPU implementation (using global memory)
- ▶ Write profiling code: Speedup

Task 2

- ▶ Make a copy of the kernel created in Task 1 and modify it to make sure that the four corner pixels are only loaded once
- ▶ Compare the performance to Task 1

Task 3

- ▶ Make a copy of the kernel created in Task 2 and use an OpenCL image for the input data
- ▶ Compare the performance to Task 1 and Task 2

OpenCL Images

- ▶ Same as CUDA “Texture Memory”
- ▶ Is a 1D / 2D / 3D array on the GPU
- ▶ Can be accessed using “samplers”
- ▶ Provide caching
 - ▶ Accesses via spatial coordinates (x,y)
- ▶ Additional features
 - ▶ Coordinate normalization
 - ▶ x/y/z coordinates go from 0.0 to 1.0
 - ▶ Return special value for out-of-bounds access
 - ▶ Filtering (i.e. linear/nearest neighbour interpolation)

OpenCL Images / Samplers

Samplers can be used to access an OpenCL Image on the GPU.

Sampler Options:

- ▶ Coordinate normalization:
 - ▶ CLK_NORMALIZED_COORDS_FALSE: Coordinates from 0 to width-1/height-1
 - ▶ CLK_NORMALIZED_COORDS_TRUE: Coordinates from 0 to 1, accessing the x coordinate by floating point value
`get_global_id(0)/get_global_size(0)`
- ▶ Addressing mode: (for out-of-bounds accesses)
 - ▶ CLK_ADDRESS_NONE: Undefined behavior
 - ▶ CLK_ADDRESS_CLAMP: Return 0
 - ▶ CLK_ADDRESS_CLAMP_TO_EDGE: Return color of border
 - ▶ CLK_ADDRESS_REPEAT: Repeat image
 - ▶ CLK_ADDRESS_MIRRORED_REPEAT: Repeat mirrored image
- ▶ Filtering:
 - ▶ CLK_FILTER_NEAREST: Nearest neighbor
 - ▶ CLK_FILTER_LINEAR: Linear/Bilinear/Trilinear interpolation

OpenCL Images / Syntax Host

Creating an Image:

```
cl::Image2D::Image2D(cl::Context context,  
    cl_mem_flags flags, cl::ImageFormat format,  
    std::size_t width, std::size_t height);
```

context = The OpenCL context to use

flags = Normally CL_MEM_READ_ONLY

format = The content of the image, e.g. cl::ImageFormat(CL_R,
CL_FLOAT), channel order CL_R = R channel, channel data type

CL_FLOAT = contains floats

width = Width of the image

height = Height of the image

OpenCL Images / Syntax Host

Copying data to an image:

```
cl::CommandQueue::enqueueWriteImage(cl::Image& image,  
    cl_bool blocking,  
    cl::size_t<3> origin, cl::size_t<3> region,  
    std::size_t row_pitch, std::size_t slice_pitch, void* ptr,  
    eventsToWaitFor = NULL, cl::Event* event = NULL) const;
```

image = The destination image

blocking = Wait until copy operation has finished (normally true)

origin = The offset in (x,y,z) in pixels in the image to write

region = The size of the region being written

row_pitch = Length of each row in bytes in the image, normally larger than (for offsets) or equal to width * sizeof(ElementType), if set to 0, the default is width * sizeof(ElementType)

slice_pitch = Bytes between two slices, for 2D images use 0

ptr = Pointer to source data

OpenCL Images / Syntax Host

Syntax for `cl::size_t<3>`:

```
cl::size_t<3> origin;  
origin[0] = origin[1] = origin[2] = 0;  
  
cl::size_t<3> region;  
region[0] = width;  
region[1] = height;  
region[2] = 1;  
queue.enqueueWriteImage(..., origin, region, ...);
```

```
Kernel.setArg<cl::Image2D>(0, image);
```

`origin` = The origin of the destination region, normally (0, 0, 0)

`region` = The size of the destination region, for 2D images normally (width, height, 1)

OpenCL Images / Syntax Kernel

Syntax for using an image:

```
// declare sampler
const sampler_t sampler = CLK_NORMALIZED_COORDS_FALSE
    | CLK_ADDRESS_CLAMP | CLK_FILTER_NEAREST;
// pass image as parameter
__kernel void sobelKernel(__read_only image2d_t d_input, ...)
// read value at (i, j)
float f = read_imagef(d_input, sampler, (int2){i, j}).x;
// write value at (i, j)
write_imagef(d_output, (int2){i,j}, (float4)x;
with x being float4 format: (grayvalue, 0, 0, 1).
```

read_imagef returns a four channel (R,G,B,A) value, where A refers to the alpha component indicating the opaqueness.

read_imagef calls that take integer coordinates must use nearest filter and unnormalized coordinates.
